



ML Study Jams

Session #3 - Classification



Vijay Jaisankar
@Vijay73893478

Tech Lead, GDSC-IIITB

```
...org = filterByOrg ? study.lead_organization === filterByOrg : true  
...status = filterByStatus ? study.status === filterByStatus : true  
...archStatus) {
```

```
function filterStudies({ studies, filterByOrg  
... filterByStatus, filterByArchStatus } = studies.filter(study
```

Classification

Essence

Categorise a set of data into a **set of classes**.

Learn the classification rules from **labelled** input data



Where do you think classification is used?

Hint: There are many 🤖




Go to slido.com and use the code
3597601

slido



Use cases of classification

① Start presenting to display the poll results on this slide.



You will “discover” a classifier
from scratch!

With the help of a few of these 
... and a few volunteers.




```
function filterStudies({ studies, filterByOrg = false, filterByStudy => {  
  studies.filter(study => {  
    ... study.organizat
```

K - Nearest Neighbours

One of the simplest classifiers


- Find the k nearest neighbours of the target point given the input data
- Assign the target point based on the majority class of the nearest neighbours
- Tie-breakers based on nearest distance





I know what you're thinking - can we
**extend Linear Regression to
classification?** We learned about it in
previous study jams and it seems
natural!

Okay, you weren't thinking about it -
but you are now!



```
function filterStudies({ studies, filterByOrg = false, filterByTopic = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organization === 'UC Berkeley';  
    }  
    if (filterByTopic) {  
      return study.topic === 'Machine Learning';  
    }  
    return true;  
  });  
}
```




Recap : Linear Regression

a.k.a “Fitting a line to data”

```
function filterStudies({ studies, filterByOrg = false, filterByYear = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organizat
```


Linear Regression - the big picture

Credits: ML Study Jams - sessions 1 and 2



Linear Regression for everyone!

a.k.a Gentle Introduction

- We estimate a target variable by representing it as a *linear sum* of input features
- The parameters of the linear sum (weights) are learned through the rinse-repeat cycle of changing them based on a *Loss function*
- At the end of training, we have a **function** that takes features' data and gives a **real number** that bears some resemblance to what the target variable would be for the same inputs.



Linear Regression - equations

Credits: ML study Jams - sessions 1 and 2

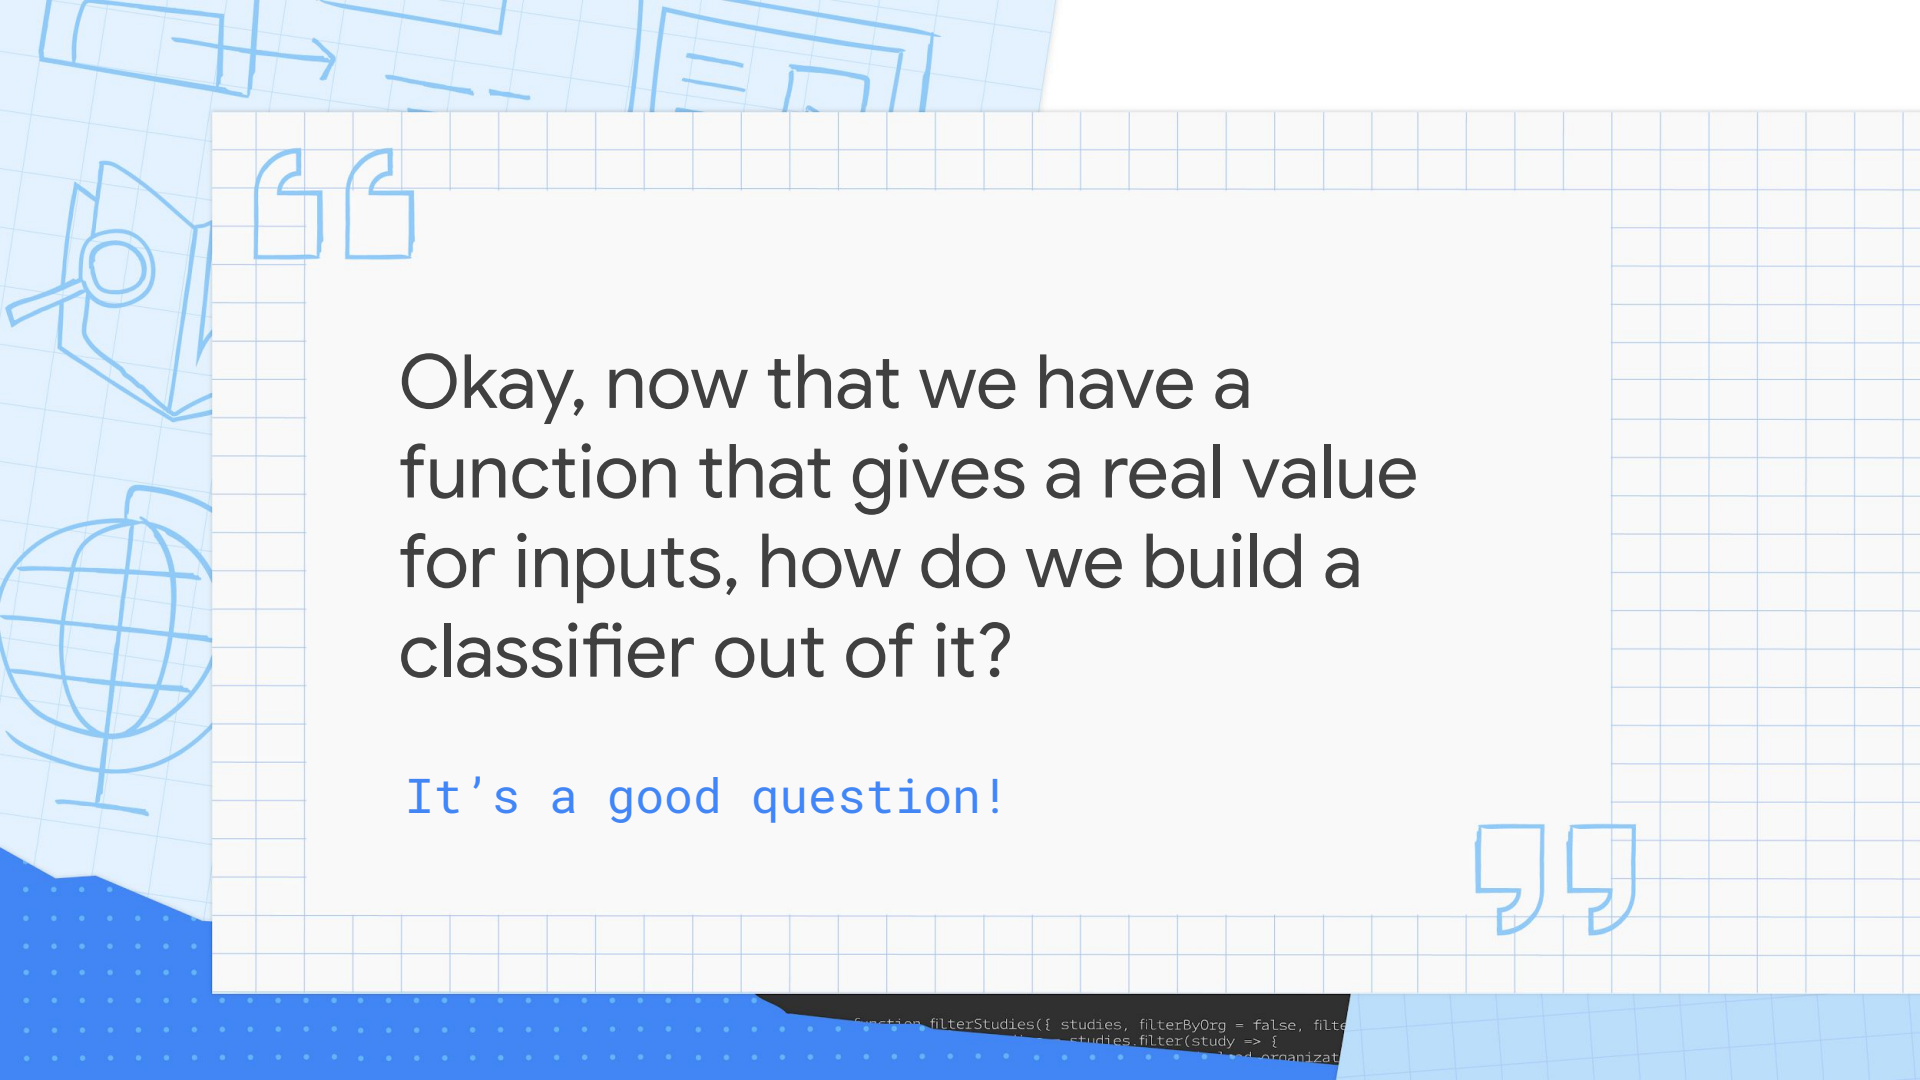
$$J(\theta) = \sum_{i=1}^N (y_i - x_i^T \theta)^2 \xrightarrow{\text{Loss function - we just saw this!}} J(\theta) = ||Y - X\theta||^2$$

$$\frac{\partial J}{\partial \theta} = 2(X^T X \theta - X^T Y)$$

Remember Gradient descent?

Updating the parameters (theta) by progressing against the gradient





Okay, now that we have a function that gives a real value for inputs, how do we build a classifier out of it?

It's a good question!

```
function filterStudies({ studies, filterByOrg = false, filterByYear = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organizat
```

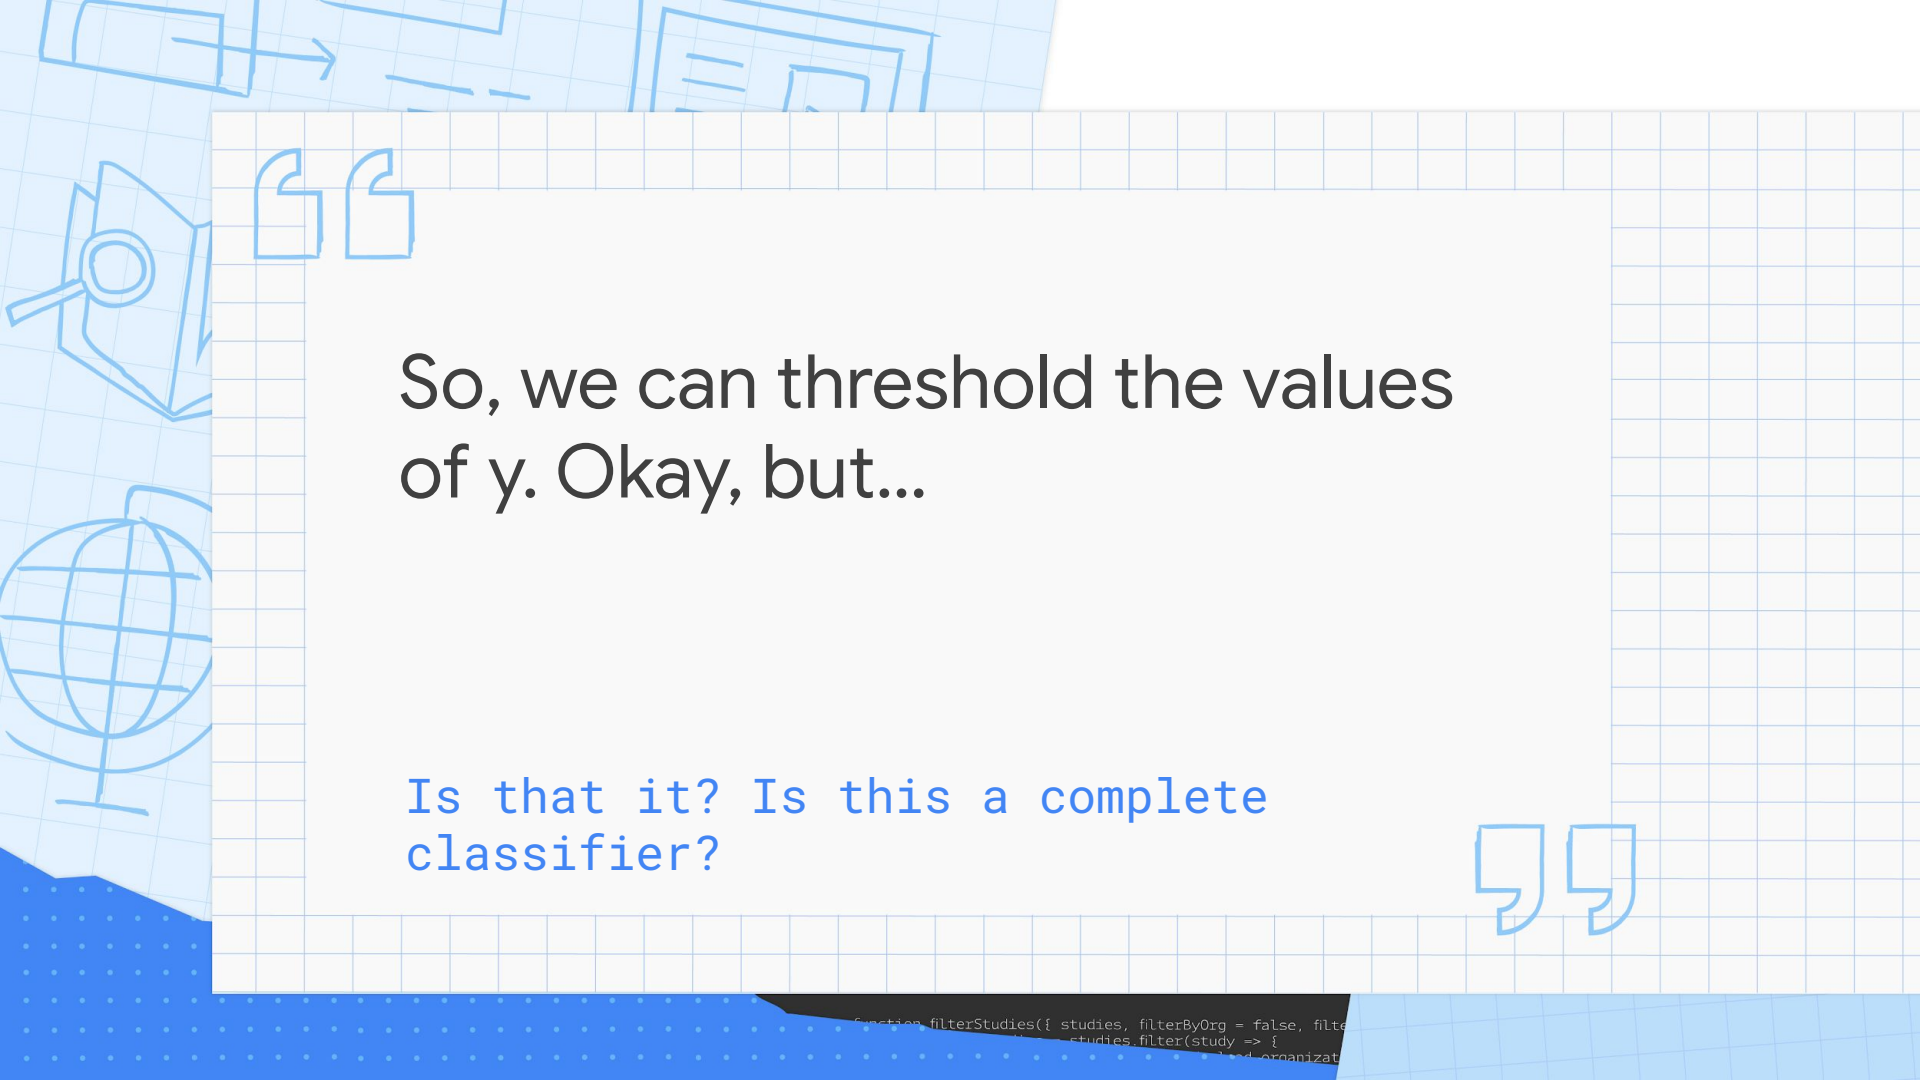


```
1  """
2      Thresholding
3          Assume binary classification
4  """
5  def condition(y):
6      pass # What will you fill here?
7
8  """
9      Find class
10         Based on the "linear sum" output, decide which class
11  """
12  def find_class_0_1(X, theta):
13      y = theta.T * X # For example
14      if condition(y):
15          return 0
16      else:
17          return 1
```



```
1  """
2      Thresholding
3          Assume binary classification
4  """
5  def condition(y):
6      pass # What will you fill here?
7
8  """
9      Find class
10         Based on the "linear sum" output, decide which class
11  """
12  def find_class_0_1(X, theta):
13      y = theta.T * X # For example
14      if condition(y):
15          return 0
16      else:
17          return 1
```

$y \leq n$



So, we can threshold the values
of y . Okay, but...

Is that it? Is this a complete
classifier?

```
function filterStudies({ studies, filterByOrg = false, filterByYear = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organizat
```

What is your intuition?

Vote now on your phones!



Stepping stones from LR to Classification

There are no right or wrong answers here, we're all learning :)



Go to slido.com and use the code
3597601

slido



Our current loss function is Mean Squared Error. Do you think that should change for Binary Classification?

① Start presenting to display the poll results on this slide.

slido




**Do you think it's easier to work with
(threshold) any value in R or a smaller
range say $[0,1]$?**

① Start presenting to display the poll results on this slide.



Logistic Regression to the rescue!

Wait, what? Logistic “Regression” is a classifier? 🙄



```
function filterStudies({ studies, filterByOrg = false, filterByYear = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organizati    }  
    if (filterByYear) {  
      return study.year > 2010    }  
    return true  })
```

First, let's look at the questions again

Remember Slido? 🗨️

- Keeping the thresholding range as **R** might come with a major unintended consequence - outliers would play a much bigger scale.
- Arithmetic errors might come into the picture, too.



Logistic Regression

The essence of the approach

- Much like Linear Regression, we learn a line.
- We feed this line into some function to get a value between 0 and 1.
- We then threshold this value to a particular class (0 or 1).



About the Loss Function

This is where things get interesting.

While RMSE is definitely valid, there's a far more effective Loss function for Binary Classification: **The Log-Loss function**


$$J = - \sum_{i=1}^N y_i \log(h_{\theta}(x_i)) + (1 - y_i) \log(1 - h_{\theta}(x_i))$$

h is the thresholded value
based on θ



How do we generate $h(\theta)$ though?

How does one convert a real number into a range from $[0,1]$?





Wow

That's a lot of maths.

```
function filterStudies({ studies, filterByOrg = false, filterByYear = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organization === 'MIT' || study.organization === 'Stanford'    }  
    if (filterByYear) {  
      return study.year > 2010  
    }  
    return true  
  })  
}
```

Logistic Regression - Summary



a.k.a “What just happened for the last 15 minutes?”

- We learn a line and threshold it using the **Sigmoid function**
- After every iteration, the loss corresponding to the current parameters are calculated through **Log loss**.
- The parameters of the line are learnt through a variation of Gradient Descent similar to Linear Regression, but the gradients will change owing to the new loss function.
- At the end of training, the output of the function will be 0/1 - a **discrete class label**!



Logistic Regression - Summary

Comparison with Linear Regression

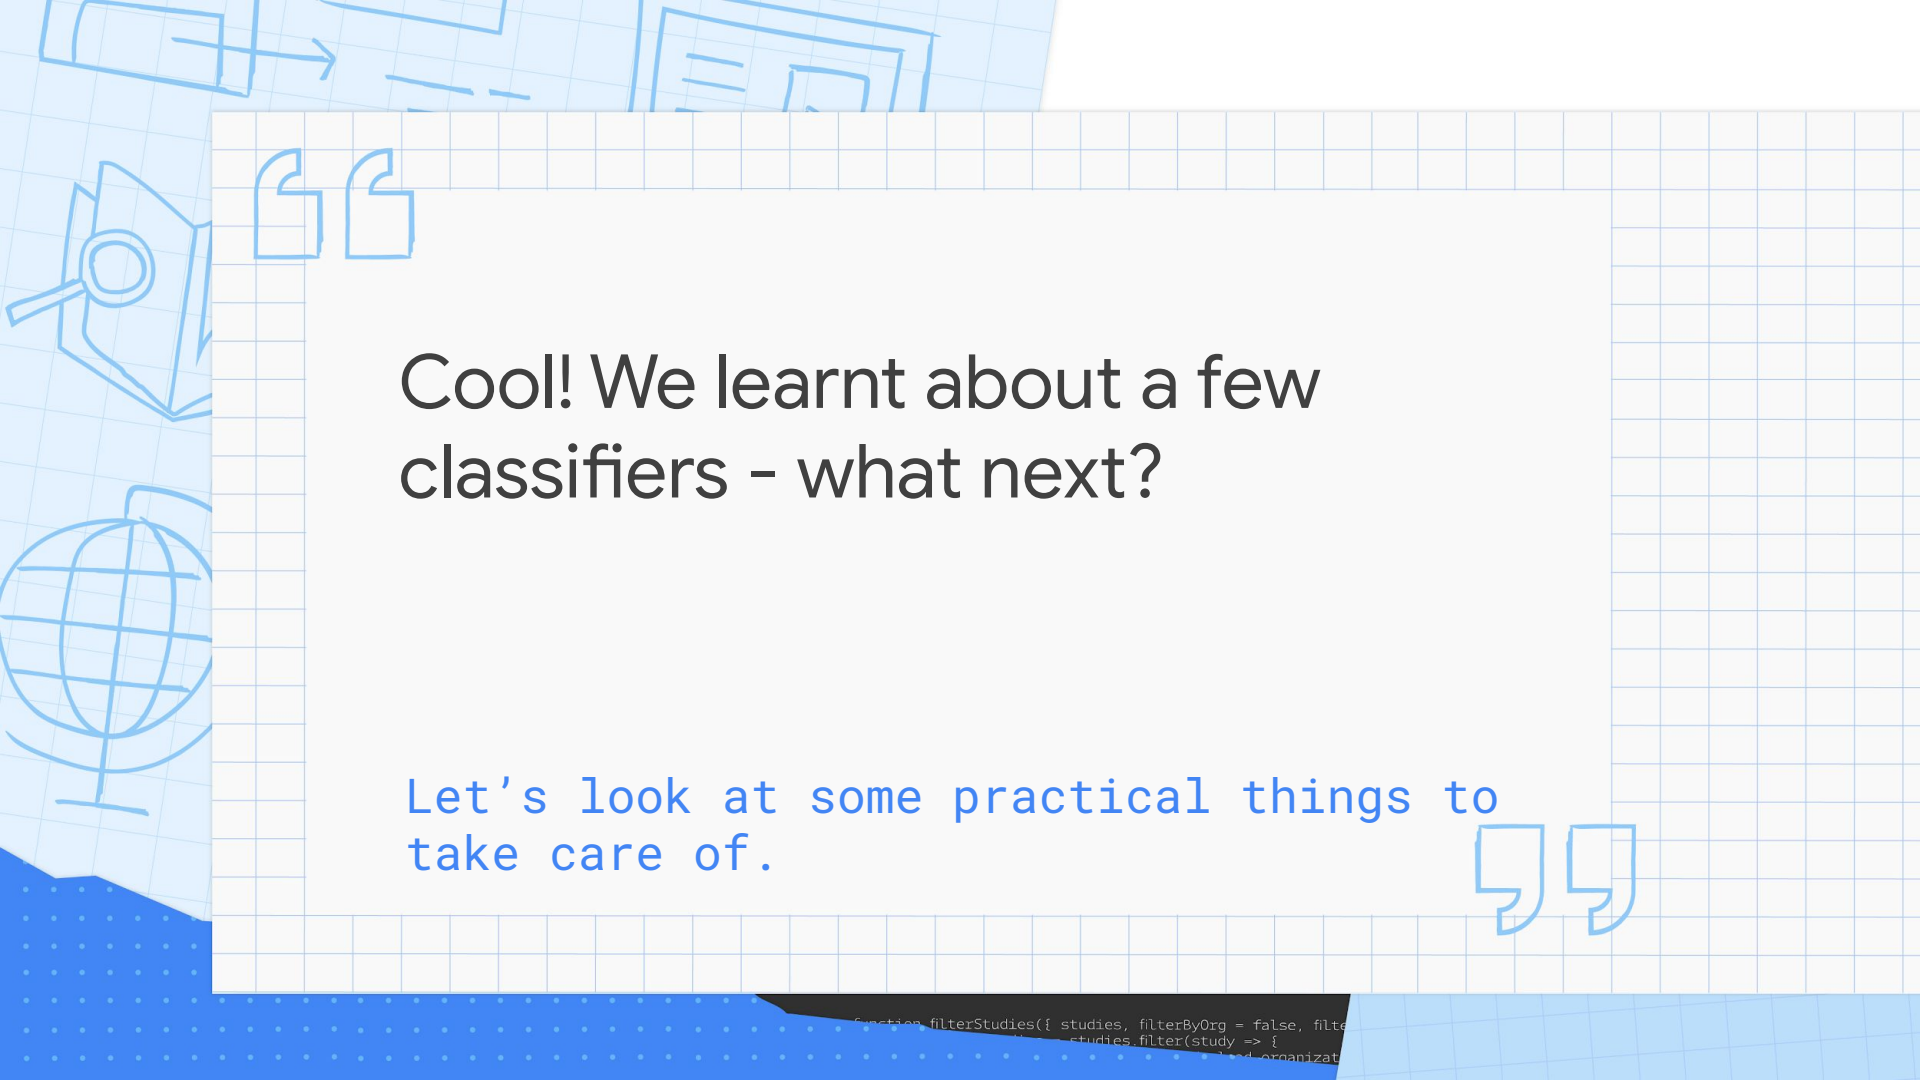


Let's build a simple LogReg Classifier!

Please make a copy of the ipynb file and let's get started!

<https://colab.research.google.com/drive/1FpeoFaWKUcq7Oeacjylq5nJdVuLBo6vm?usp=sharing>





Cool! We learnt about a few classifiers - what next?

Let's look at some practical things to take care of.

```
function filterStudies({ studies, filterByOrg = false, filterByTopic = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organizat
```

PSA

From now on, consider your classifier to be a “black box”.

Dealing with multiple classes

More than just 0/1

- One vs All method
 - Make k classifiers - one for each class
 - Most confident class
- Some classifiers have support for multiple classes
 - More complex loss functions and prediction functions
 - Look into Multinomial Log Loss

Dealing with Class Imbalance

Relative frequency of data points

- Oversampling minority class
 - Example: SMOTE
- Undersampling majority class
- Passing `class_weights` as a parameter to the model

Dealing with Class Imbalance

How about testing the predictions?

This is a loaded question.



Fraud Detection Problem - ~ 200 features, ~120000 rows.

Binary Classification - 0 for no fraud
and 1 for fraud

“

A “model” trained in under 2 minutes and had a 95% accuracy!

Can you guess what it was?

”

```
function filterStudies({ studies, filterByOrg = false, filterByYear = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organization !== 'unlabeled-organization'  
    }  
    if (filterByYear) {  
      return study.year !== 'unlabeled-year'  
    }  
    return true  
  })  
}
```



```
1  """
2      Crazy Model
3          Trained under 2 minutes and gave 95% accuracy
4  """
5  def get_predictions(X_test):
6      # What do you think goes here?
```



```
1  """
2      Crazy Model
3          Trained under 2 minutes and gave 95% accuracy
4  """
5  def get_predictions(X_test):
6      return [0] * len(X_test)
```

**It just returned "Not fraud"
for all samples!**

Dealing with Class Imbalance


Sometimes, accuracy is not enough

We need to account for class frequencies too. A good tester is the **F1 score** metric that weights false positives and false negatives too!

Dealing with underfitting and overfitting

a.k.a “nah m8 not happening” and “suffering from success”

- Underfitting
 - Increasing training time
 - Making the hypothesis more complex
 - Making the data more streamlined
- Overfitting
 - Validation sets and validation testing
 - Cross-validation



“

There are many classifiers.

Logistic regression

Naive Bayes

KNN


SVM

The list goes on and on...

”



```
function filterStudies({ studies, filterByOrg = false, filterByYear = false }) {
  return studies.filter(study => {
    if (filterByOrg) {
      return !study.organization;
    }
    if (filterByYear) {
      return !study.year;
    }
    return true;
  });
}
```

“

So, how would you choose the best classifier for your problem?

Try them out one-by-one and choose the best classifier based on a testing metric.

”




```
function filterStudies({ studies, filterByOrg = false, filterByYear = false, filterByOrganizati
```

Can we do better?

The PTSD associated with this statement...

Should we only stick to **one** model?

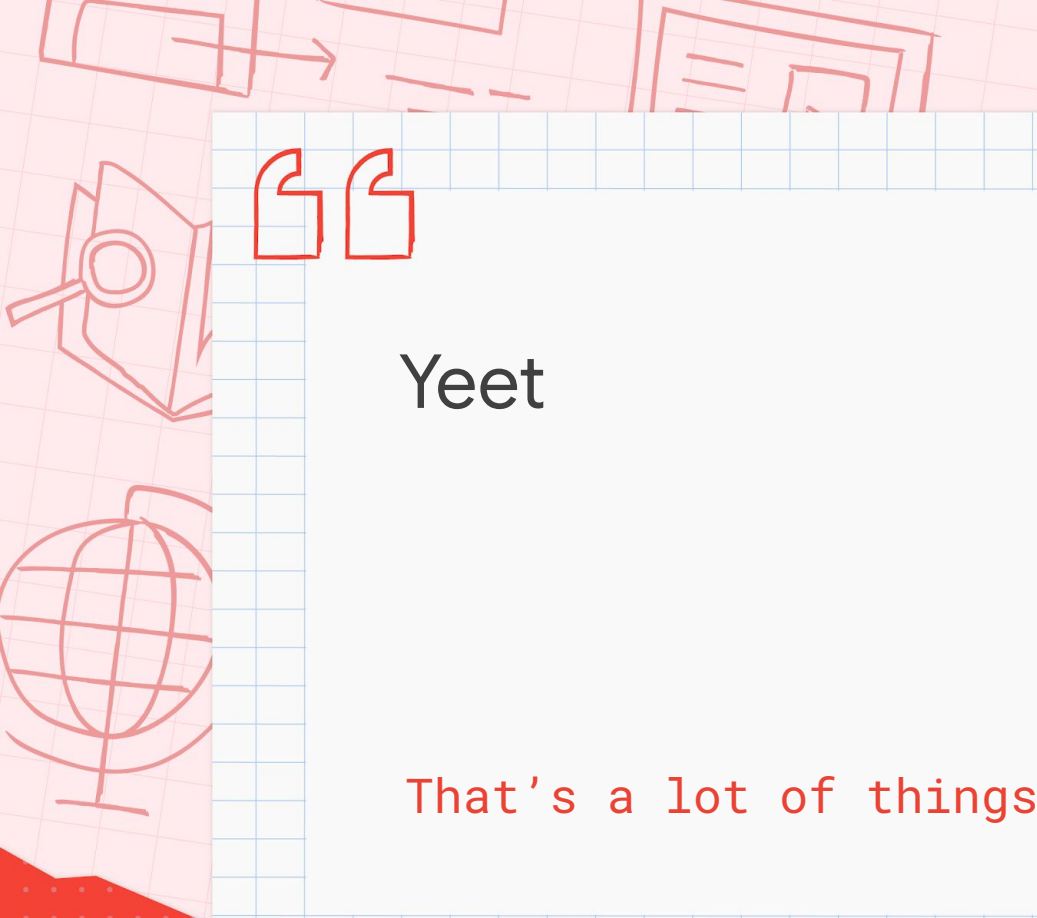
What if we can combine models? Won't that be good?



Ensemble Methods

Combining multiple models

- Bagging
 - Take m models
 - Find out the predictions of each model
 - Take the most frequent prediction
 - Does this sound familiar? 🤖
- There's also boosting and stacking, which is left as an exercise to the reader 🤖



“

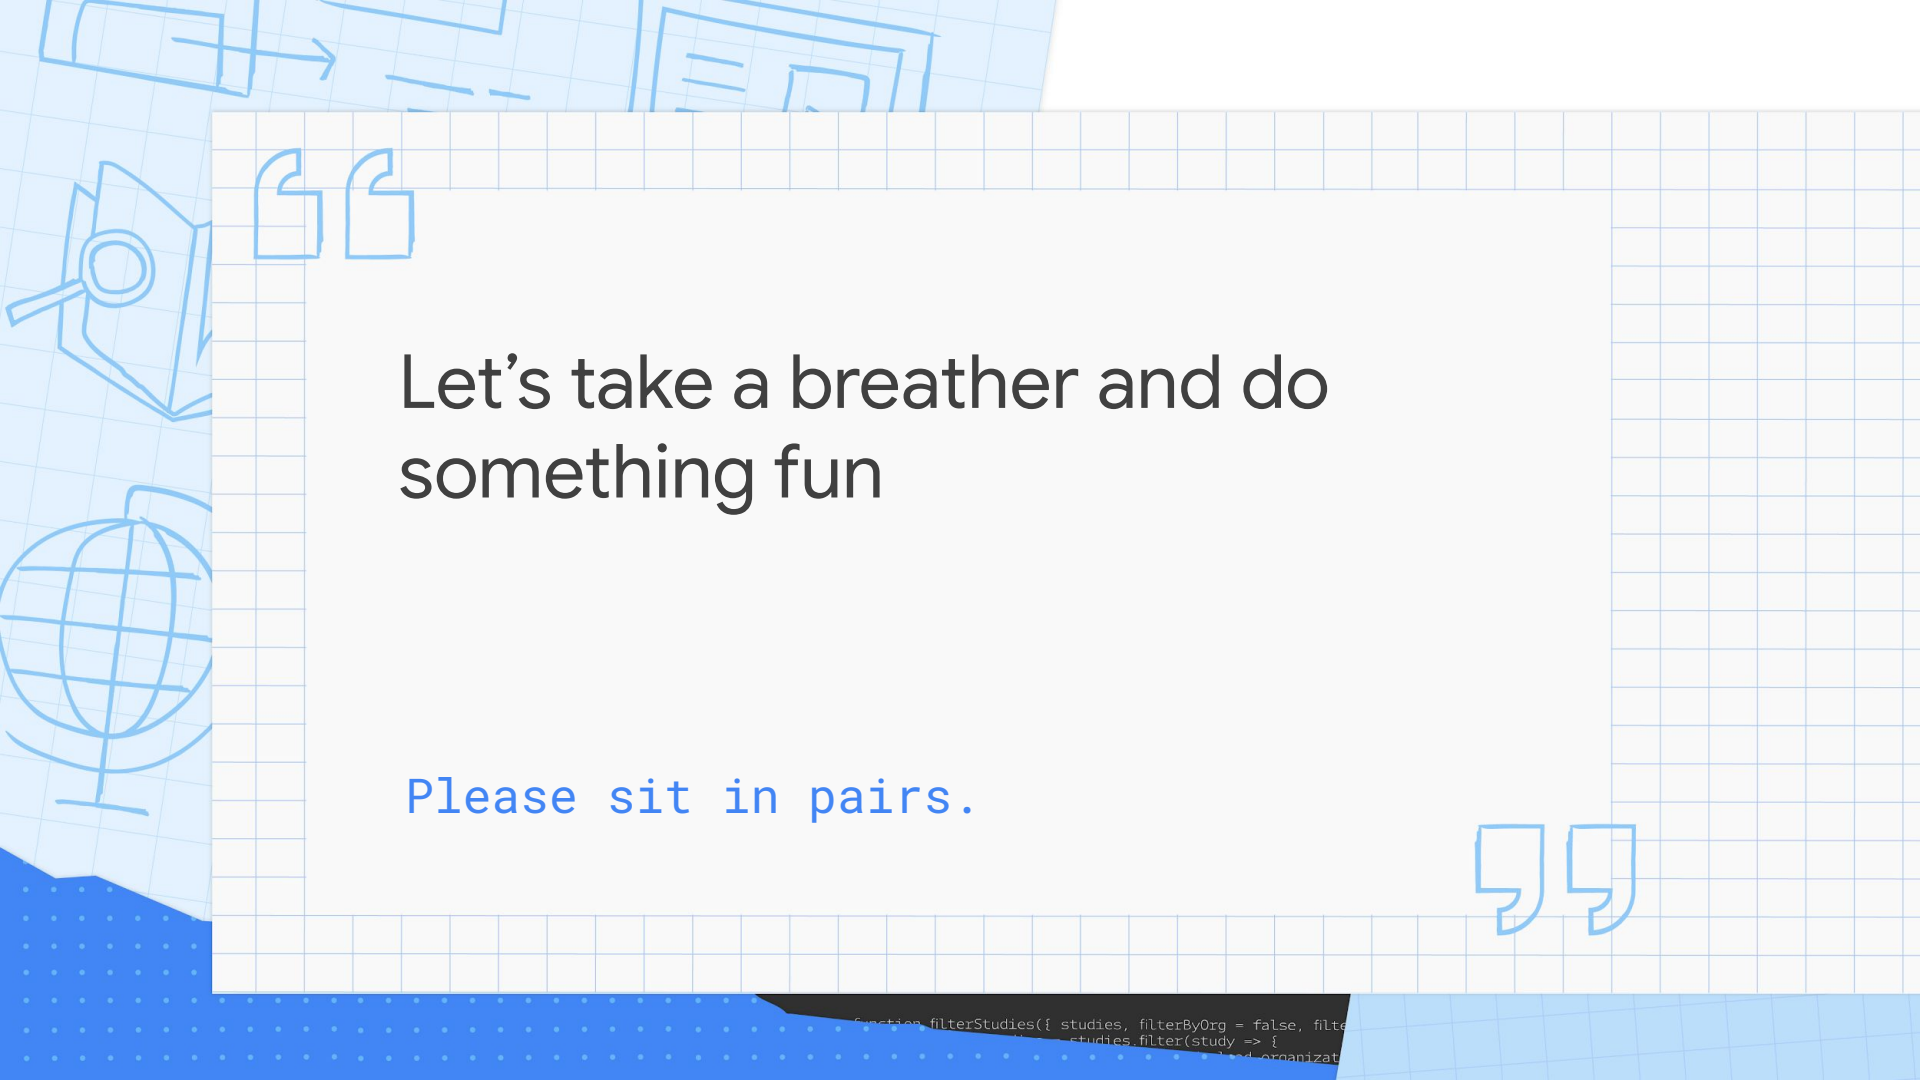
Yeet

That's a lot of things to take care of.

”



```
function filterStudies({ studies, filterByOrg = false, filterByYear = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organization === 'NIH'  
    }  
    if (filterByYear) {  
      return study.year === 2020  
    }  
    return true  
  })  
}
```



Let's take a breather and do
something fun

Please sit in pairs.

```
function filterStudies({ studies, filterByOrg = false, filterByTopic = false }) {  
  return studies.filter(study => {  
    if (filterByOrg) {  
      return study.organizat
```

Teachable Machine

Note: You need to enable your webcams for this.



<https://teachablemachine.withgoogle.com/train/image>

Teachable Machine

How do they do it?

- Very complex neural networks!
- Transfer learning



Does that sound interesting to you?

Please say yes :)

Find out more about these topics in our next session!

Session #4: **Big-brain time with neural networks**



Thank you!

If you'd like to post about the event,
please tag [@gdsc_iitb](#)



Any questions?

